

Building Multi-Linear Regression from Scratch

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_california_housing
from sklearn.preprocessing import StandardScaler
```

Importing the libraries

```
# Load dataset
data = fetch_california_housing()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['Target'] = data.target
print(df)
# Split dataset
x = df.drop('Target', axis=1)
y = df['Target']
```

20640 Samples

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	Target
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422
...
20635	1.5603	25.0	5.045455	1.133333	845.0	2.560606	39.48	-121.09	0.781
20636	2.5568	18.0	6.114035	1.315789	356.0	3.122807	39.49	-121.21	0.771
20637	1.7000	17.0	5.205543	1.120092	1007.0	2.325635	39.43	-121.22	0.923
20638	1.8672	18.0	5.329513	1.171920	741.0	2.123209	39.43	-121.32	0.847
20639	2.3886	16.0	5.254717	1.162264	1387.0	2.616981	39.37	-121.24	0.894

[20640 rows x 9 columns]

8 Features,
x_size: 20640 x 8

Target,
y_size: 20640 x 1

```
train_split = 0.8
train_split_index = int(train_split * len(x))
```

```
x_train = x.iloc[:train_split_index]
x_test = x.iloc[train_split_index:]
y_train = y.iloc[:train_split_index]
y_test = y.iloc[train_split_index:]
```

```
# Scale data
scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
print(' x_train_shape:', x_train.shape, '\n', 'x_test_shape:', x_test.shape, '\n', 'y_train_shape:', y_train.shape, '\n', 'y_test_shape:', y_test.shape)
```

Note: Without standardization, the gradient descent updates can become unstable, leading to exploding or vanishing gradients. This happens because features with large magnitudes dominate the weight updates, causing inefficient learning and convergence issues.


```
x_train_shape: (16512, 8)
x_test_shape: (4128, 8)
y_train_shape: (16512,)
y_test_shape: (4128,)
```

```
x_train_size: (0.8 x 20640) x 8
x_train_size: 16512 x 8
x_test_size: (0.2 x 20640) x 8
x_test_size: 4128 x 8
```


Since y represents the target variable, it has a single column

```
y_train_size: 16512 x 1
y_test_size: 4128 x 1
```


Selecting Learning Rates



It represents a scenario where the learning rate is low, allowing the optimization process to carefully explore the loss landscape and potentially converge to a local minimum.



It shows a higher learning rate, which can cause the optimizer to take large steps, potentially skipping over global minima and moving toward another region.



Adjusting the learning rate properly ensures the optimizer reaches the weight corresponding to the minimum cost function (global minimum)

```
class Multi_Linear_Regression:
    def __init__(self, lr=0.01, n_iter=1000):
        self.lr = lr
        self.n_iter = n_iter

    def fit(self, X_train, y_train):
        n_samples, n_features = X_train.shape
        self.weights = np.zeros(n_features)
        self.bias = 0

        for _ in range(self.n_iter):
            y_pred = np.dot(X_train, self.weights) + self.bias
            dw = (1/n_samples) * np.dot(X_train.T, (y_pred - y_train))
            db = (1/n_samples) * np.sum(y_pred - y_train)
            self.weights -= self.lr * dw
            self.bias -= self.lr * db

    def predict(self, X_test):
        return np.dot(X_test, self.weights) + self.bias

    def evaluate(self, X_test, y_test):
        y_pred = self.predict(X_test)
        r_square = 1 - (np.sum((y_test - y_pred) ** 2) / np.sum((y_test - np.mean(y_test)) ** 2))
        print(f"R2 Score: {r_square:.4f}")
        return r_square
```

Prediction Equation: $\hat{y} = X_{\text{train}} \cdot w + b$

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_{16512} \end{bmatrix} = \begin{bmatrix} x_{1,1} & x_{1,2} & \dots & x_{1,8} \\ x_{2,1} & x_{2,2} & \dots & x_{2,8} \\ \vdots & \vdots & \ddots & \vdots \\ x_{16512,1} & x_{16512,2} & \dots & x_{16512,8} \end{bmatrix} \cdot \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_8 \end{bmatrix} + b$$

w and b are initialized as 0, and is updated using gradient descent.

Gradient Descent Rule:

$$w = w - lr \cdot dw$$

$$b = b - lr \cdot db$$

Cost function (Mean Square Error):

$$J(w, b) = \frac{1}{2n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

Substituting \hat{y} into the Cost Function

$$J(w, b) = \frac{1}{2n} \sum_{i=1}^n (y_i - (X_i \cdot w + b))^2$$

Differentiate $J(w, b)$ w.r.t. w :

$$\frac{dJ}{dw} = \frac{1}{2n} \cdot 2 \sum_{i=1}^n (y_i - (X_i \cdot w + b)) \cdot (-X_i)$$

$$\frac{dJ}{dw} = \frac{1}{n} \sum_{i=1}^n X_i \cdot (-y_i - (X_i \cdot w + b))$$

$$\frac{dJ}{dw} = \frac{1}{n} \sum_{i=1}^n X_i \cdot ((X_i \cdot w + b) - y_i)$$

$$dw = \frac{1}{n} X^T (\hat{y} - y)$$

Differentiate $J(w, b)$ w.r.t. b :

$$\frac{dJ}{db} = \frac{1}{2n} \cdot 2 \sum_{i=1}^n (y_i - (X_i \cdot w + b)) \cdot (-1)$$

$$\frac{dJ}{db} = \frac{1}{n} \sum_{i=1}^n ((X_i \cdot w + b) - y_i)$$

$$db = \frac{1}{n} \sum_{i=1}^n (\hat{y} - y)$$

Equation for R^2

$$R^2 = 1 - \frac{\sum (y_{\text{test}} - \hat{y})^2}{\sum (y_{\text{test}} - \bar{y})^2}$$

```
# Train Multi-Linear Regression Model
model = Multi_Linear_Regression(lr=0.01, n_iter=1000)
model.fit(x_train_scaled, y_train.values)

# Evaluate the model
model.evaluate(x_test_scaled, y_test.values)
```

Multi-Linear Regression Model using gradient descent with a learning rate of 0.01 and 1000 iterations is initialized. After training on the scaled dataset, the model is evaluated on the test data to compute the R^2 score, measuring its performance.